

Week 01 - Data Pre-Processing

Aims

By the end of this notebook you will

- understand and play with the different aspects of data pre-processing
- be familiar with tools for exploratory data analysis and visualization
- understand the basics of feature engineering

Topics and Instructions

1. [Problem Definition and Setup](#)
2. [Exploratory Data Analysis](#)
3. [Data Preprocessing](#)
4. [Feature Engineering](#)
5. [Summary](#)

In lecture this week, we reviewed the general **machine learning pipeline**, which following the "[Machine Learning Project Checklist](https://github.com/ageron/handson-ml/blob/master/ml-project-checklist.md)" (<https://github.com/ageron/handson-ml/blob/master/ml-project-checklist.md>) of Geron (2019) can be structured as:

- Frame the problem and look at the big picture.
- Get the data.
- Explore the data and gain insights.
- Prepare the data to better expose the underlying data patterns to Machine Learning algorithms.
- Explore many different models and shortlist the best ones.
- Fine-tune your models and combine them into a great solution.
- Present your solution.
- Launch, monitor, and maintain your system.

In this week's workshop, we will focus on the initial steps of this pipeline, that is on, data pre-processing, exploratory data analysis and feature engineering.

During workshops, you will complete the worksheets together in teams of 2-3, using **pair programming**. During the first few weeks, the worksheets will contain cues to switch roles between driver and navigator. When completing worksheets:

- You will have tasks tagged by (CORE) and (EXTRA).
- Your primary aim is to complete the (CORE) components during the WS session, afterwards you can try to complete the (EXTRA) tasks for your self-learning process.
- Look for the 🔄 as cue to switch roles between driver and navigator.
- In some Exercises, you will see some beneficial hints at the bottom of questions.

Instructions for submitting your workshops can be found at the end of worksheet. As a reminder, you must submit a pdf of your notebook on Learn by 16:00 PM on the Friday of the week the workshop was given.

Problem Definition and Setup

Packages

Now lets load in some packages to get us started. The following are widely used libraries to start working with Python in general.

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import seaborn as sns
        4 import matplotlib.pyplot as plt
        5 from scipy import stats
```

C:\Users\TRISHA\anaconda3\lib\site-packages\pandas\core\computation\expressions.py:21: UserWarning: Pandas requires version '2.8.4' or newer of 'numexpr' (version '2.8.3' currently installed).

from pandas.core.computation.check import NUMEXPR_INSTALLED

C:\Users\TRISHA\anaconda3\lib\site-packages\pandas\core\arrays\masked.py:60: UserWarning: Pandas requires version '1.3.6' or newer of 'bottleneck' (version '1.3.5' currently installed).

from pandas.core import (

If you need to install any packages from scratch, you need to install the related library before calling it. For instance, [feature-engine](https://feature-engine.trainindata.com/en/latest/) (<https://feature-engine.trainindata.com/en/latest/>) is a Python library for Feature Engineering and Selection, which:

- contains multiple transformers to engineer and select features to use in machine learning models.
- preserves scikit-learn functionality with methods `fit()` and `transform()` to learn parameters from and then transform the data.

```
In [2]: 1 # To install the feature-engine library (if not already installed)
        2 # !pip install feature-engine
```

In some cases, we may need only a component of the whole library. If this is the case, it is possible to import specific things from a module (library), using the following line of code:

```
In [3]: 1 from feature_engine.imputation import DropMissingData
```

Problem

Now, it is time move on to the next step.

Welcome to Machine Learning Housing Corporation! The first task you are asked to perform is to build a model of housing prices in California using the California census data. This data has metrics such as the population, median income, median housing price, and so on for each block group in California. Block groups are the smallest geographical unit for which the US Census Bureau publishes sample data (a block group typically has a population of 600 to 3,000 people). We will just call them “districts” for short.

Your model should learn from this data and be able to predict the median housing price in any district, given all the other metrics.

The first question to ask your boss is what exactly is the business objective; building a model is probably not the end goal. **How does the company expect to use and benefit from this model?** This is important because it will determine how you frame the problem, what algorithms you will select, what performance measure you will use to evaluate your model, and how much effort you should spend tweaking it.

The next question to ask is **what the current solution looks like (if any)**. It will often give you a reference performance, as well as insights on how to solve the problem. Your boss answers that the district housing prices are currently estimated manually by experts: a team gathers up-to-date information about a district, and when they cannot get the median housing price, they estimate it using complex rules.

This is costly and time-consuming, and their estimates are not great; in cases where they manage to find out the actual median housing price, they often realize that their estimates were off by more than 20%. This is why the company thinks that it would be useful to train a model to predict a district’s median housing price given other data about that district. The census data looks like a great dataset to exploit for this purpose, since it includes the median housing prices of thousands of districts, as well as other data.

► Exercise 1 (CORE)

Using the information above answer the following questions about how you may design your machine learning system.

a) Is this a supervised or unsupervised learning task?

Supervised

b) Is this a classification, regression, or some other task?

Regression

c) Suppose you are only required to predict if a district's median housing prices are "cheap," "medium," or "expensive". Will this be the same or a different task?

No, it would become a classification task

Data Download

The data we will be using this week is a modified version of the California Housing dataset. We can get the data a number of ways. The easiest is just to load it from the working directory that we are working on (where we have already downloaded it to).

```
In [4]: 1 housing = pd.read_csv("housing.csv")
2 housing.index = range(len(housing))
3
4 housing.head()
5
```

```
Out[4]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_proximity
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	452600.0	near bay or ocean
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	358500.0	near bay or ocean
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	352100.0	near bay or ocean
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	341300.0	near bay or ocean
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	342200.0	near bay or ocean

Exploratory Data Analysis

In this section we are going to start with exploring the California Housing data using methods that you will likely already be familiar with.

Data can come in a broad range of forms encompassing a collection of discrete objects, numbers, words, events, facts, measurements, observations, or even descriptions of things. Processing data using exploratory data analysis (EDA) can elicit useful information and knowledge by examining the available dataset to discover patterns, spot anomalies, test hypotheses, and check assumptions.

Let's start by examining the [Data Dictionary \(https://www.kaggle.com/camnugent/california-housing-prices\)](https://www.kaggle.com/camnugent/california-housing-prices) and the variables available:

```
longitude : A measure of how far west a house is; a higher value is farther west

latitude : A measure of how far north a house is; a higher value is farther north

housingMedianAge : Median age of a house within a block; a lower number is a newer building

totalRooms : Total number of rooms within a block

totalBedrooms : Total number of bedrooms within a block

population : Total number of people residing within a block

households : Total number of households, a group of people residing within a home unit, for a block

medianIncome : Median income for households within a block of houses (measured in tens of thousands of US Dollars)

medianHouseValue : Median house value for households within a block (measured in US Dollars)

oceanProximity : Location of the house w.r.t ocean/sea
```

► Exercise 2 (CORE)

a) Examine the datatypes for each column calling `info()`. (<https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.info.html>). What is the total number of observations and total number of variables? What is the type of each variable?

```
In [5]: 1 housing.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population             20640 non-null  float64
6   households             20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

Number of observations : 20640

No of variables: 10

b) From the information provided above, can you anticipate any data cleaning we may need to do?

Type *Markdown* and LaTeX: α^2

Exercise 3 (CORE)

a) Use descriptive statistics and histograms to examine the distributions of the numerical attributes.


Hint

In [6]: 1 housing.describe()

Out[6]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000	20640.000000	20640.000000	
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.476744	499.539680	3.870671	:
std	2.003532	2.135952	12.585558	2181.615252	421.385070	1132.462122	382.329753	1.899822	
min	-124.350000	32.540000	1.000000	2.000000	1.000000	3.000000	1.000000	0.499900	
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.000000	280.000000	2.563400	
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.000000	409.000000	3.534800	.
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.000000	605.000000	4.743250	:
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000	6082.000000	15.000100	!

```
In [7]: 1 plt.figure(figsize=(8, 5))
2 sns.histplot(housing['housing_median_age'], bins=50, kde=False, color='skyblue', edgecolor='black')
3 plt.show()
4 sns.histplot(housing['latitude'], bins=50, kde=False, color='skyblue', edgecolor='black')
5 plt.show()
6 sns.histplot(housing['longitude'], bins=50, kde=False, color='skyblue', edgecolor='black')
7 plt.show()
8 sns.histplot(housing['total_rooms'], bins=50, kde=False, color='skyblue', edgecolor='black')
9 plt.show()
10 sns.histplot(housing['total_bedrooms'], bins=50, kde=False, color='skyblue', edgecolor='black')
11 plt.show()
12 sns.histplot(housing['population'], bins=50, kde=False, color='skyblue', edgecolor='black')
13 plt.show()
14 sns.histplot(housing['households'], bins=50, kde=False, color='skyblue', edgecolor='black')
15 plt.show()
16 sns.histplot(housing['median_income'], bins=50, kde=False, color='skyblue', edgecolor='black')
17 plt.show()
18 sns.histplot(housing['median_house_value'], bins=50, kde=False, color='skyblue', edgecolor='black')
19 plt.show()
20
21 #could have just used a loop after attesting the information about the columns first row.
```



b) Can you identify other pre-processing/feature engineering steps we may need to do?

Type *Markdown* and LaTeX: α^2

c) One thing you may have noticed from the histogram is that the median income, housing median age, and the median house value are capped. The median house value capping (this being our target value) may or may not be a problem depending on your client. If we needed precise predictions beyond \$500,000, we may need to either collect proper labels/outputs for the districts whose labels were capped or

remove these districts from the data. Following the latter, remove all districts whose median house value is capped. How many observations are there now?

```
In [8]: 1 # Remove the cases where median_house_value >= 500,000$
        2 housing = housing[housing["median_house_value"] < 500000]
        3
```

► Exercise 4 (CORE)


What are the possible categories for the `ocean_proximity` variable? Are the number of instances in each category similar?

Hint

- `value_counts()` (https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.Series.value_counts.html) can be used to count the values of the categories a pandas series.
- You can use a `sns.countplot` (<https://seaborn.pydata.org/generated/seaborn.countplot.html>) to create barplot with the number of instances of each category

```
In [9]: 1 housing["ocean_proximity"].value_counts()
        2
        3 sns.countplot(housing["ocean_proximity"])
```

Out[9]: <AxesSubplot:xlabel='count', ylabel='ocean_proximity'>



No, the number of instances in each category are not the same

⚠ Now, is a good point to switch driver and navigator

► Exercise 5 (CORE)


Examine if/which of the features are correlated to each other. Are any of the features correlated with our output (`median_house_value`) variable?

- Can you think of any reason why certain features may be correlated?
- How might we use this information in later steps of our model pipeline?

Hint

- `.corr()` can be used to compute the correlations.
- You can use a `sns.heatmap` (<https://seaborn.pydata.org/generated/seaborn.heatmap.html>) to visualize the correlations

```
In [10]: 1 #got an error stating could not convert string to float: 'NEAR BAY', work on this.
2 housing_numeric = housing.select_dtypes(include=['float64', 'int64'])
3 correlation_matrix = housing_numeric.corr()
4 plt.figure(figsize=(10, 10))
5 sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm", fmt=".2f", linewidths=0.5)
6 plt.title("Correlation Matrix")
7 plt.show()
```



Median Income Higher median income is strongly associated with higher house prices, as wealthier areas usually have more expensive homes.

Weak Positive Correlations: Older housing areas and a greater number of rooms show weak positive correlations with house prices.

Weak Negative Correlations: Moving northward in California slightly decreases house prices, while moving eastward has a very weak negative impact on prices.

Negligible Correlations: Features like population, households, and total bedrooms have minimal correlations with house values.


Are any features correlated with each other? Yes, total rooms, total bedrooms, population, and households are highly correlated, reflecting similar aspects of housing density or size.

Why might certain features be correlated? Wealthy areas have higher demand for housing, driving up prices. Housing density measures and geographical factors also influence house prices. ch impacts house prices.

► Exercise 6 (CORE)

Use `sns.pairplot` to further investigate the joint relationship between each pair of variables. What insights into the data might this provide over looking only at the correlation?

```
In [11]: 1 selected_features = ["median_income", "housing_median_age", "total_rooms",
2                       "population", "median_house_value"]
3
4 #sns.pairplot(housing_numeric[selected_features], diag_kind="kde", corner=True)
5 sns.pairplot(housing[selected_features], diag_kind="kde", corner=True)
6 plt.show()
```



Median Income vs. Median House Value: There is a strong positive relationship, indicating that higher incomes are associated with higher house prices.


Total Rooms/Population vs. Median House Value: A weak or unclear relationship is observed here. This suggests that these features may need transformation, such as using ratios like rooms per household.

Housing Median Age vs. Median House Value: There is a slight positive trend, suggesting that older neighborhoods may have higher house values.

► Exercise 7 (CORE)

Which variables in our data represent counts and how are they distributed? Use a probability plot `scipy.stats.probplot(df[variable], dist="norm")` to compare their distributions with a normal.

```
In [18]: 1 from scipy.stats import probplot
2
3 count_variables = ["total_rooms", "total_bedrooms", "population", "households"]
4
5 for variable in count_variables:
6     plt.figure(figsize=(8, 5))
7     probplot(housing[variable], dist="norm", plot=plt)
8     plt.title(f"Probability Plot of {variable} vs Normal Distribution")
9     plt.grid(axis='both', linestyle='--', alpha=0.7)
10    plt.tight_layout()
11    plt.show()
12
```



Skewness: All count variables show significant deviations from the norm, as indicated by the red line, suggesting they are not normally distributed. There is a pronounced right skew, evident from the long tails at higher values.

Outliers: Extreme values (outliers) at the higher end contribute to the deviations from normality.

Transformations: Applying transformations, such as logarithmic or square root, may help reduce skewness and bring these distributions closer to a normal shape.

Data Pre-Processing

Now we have some familiarity with the data though EDA, let's start preparing our data to be modelled.

From here on in, we should really be creating functions for our data transformations. This is because when we want to run data through our "model pipeline" in the future, rather than having to copy and paste a load of code, we can just use a series of functions. Later on the course, we will see this is useful when we split our data into training, validation, and test sets, but this would also be required if you deploy your model in a "live" environment. Furthermore, when refining a model it makes it easier for us to treat our preparation choices as "hyperparameters", meaning we can easily add or remove parts of our pipeline to see what works and what doesn't.

It's also worth examining what is meant by a **"Pipeline"**.

- A general definition is that it is just a sequence of data preparation operations that is ensured to be reproducible.
- However, we may want to ensure that any functions/classes we make for our pipeline have specific attributes that work best with the tools available in our chosen machine learning library.
- In this course we are mostly going to be using `Scikit-learn`, with a little `Keras` at the end for neural networks.

Scikit-learn

Scikit-learn is an open source machine learning library that supports supervised and unsupervised learning (https://scikit-learn.org/stable/getting_started.html (https://scikit-learn.org/stable/getting_started.html)).

It also provides various tools for model fitting, data preprocessing, model selection and evaluation, and many other utilities.

In Scikit-Learn a `Pipeline` is a class we can use to combine our pre-processing and modelling steps together (<https://scikit-learn.org/stable/modules/generated/sklearn.pipeline.Pipeline.html> (<https://scikit-learn.org/stable/modules/generated/sklearn.pipeline.Pipeline.html>)).

A `Pipeline` is useful for many reasons, including that they help prevent you from data leakage, i.e. disclosing some testing data in your training data. Objects that go into a scikit-learn `Pipeline` can either be *transformer* or *estimator* classes, or, if we use an `imbalanced-learn` `Pipeline` instead, also *resamplers*.

All three of these objects (*resamplers*, *transformers*, and *estimator*) all typically have a `.fit()` method, which is used to

- validate and interpret any parameters,
- validate the input data,
- estimate and store attributes from the parameters and provided data,

- return the fitted estimator to facilitate method chaining in a pipeline.

Along with other sample properties (e.g. `sample_weight`), the `.fit()` method usually takes two inputs:

- The input matrix (or design matrix) `X`. The size of `X` is typically `(n_samples, n_features)`, which means that samples are represented as rows and features are represented as columns.
- The target values `y` which are real numbers for regression tasks, or integers for classification (or any other discrete set of values). For unsupervised learning tasks, `y` does not need to be specified.

https://scikit-learn.org/stable/getting_started.html (https://scikit-learn.org/stable/getting_started.html)

Other methods available for these objects other than `.fit()` will depend on what they are, so we will look at each of these objects in turn and then combine them into a model pipeline later in this workshop.

In general, it is good practice to adhere the **Scikit-Learn Design principles** when creating our pipeline steps.

- **Nonproliferation of classes.** Datasets are represented as NumPy arrays or SciPy sparse matrices, instead of homemade classes. Hyperparameters are just regular Python strings or numbers.
- This means all our functions should **input and output Numpy arrays**. You can of course use pandas functions if you wish in your function, but just make sure that you convert the data back to output the Numpy array object!
- Otherwise, in some steps, there might appear some object/data type errors

Lets start by splitting our **features** from our **target** variable in the data set.

```
In [12]: 1 # Extracting the features from the data
2 X = housing.drop("median_house_value", axis = 1)
3 features = list(X.columns)
4 print(features)
5
6 X = X.values
7 print(X) # Or display(X)
8 print(X.shape)

['longitude', 'latitude', 'housing_median_age', 'total_rooms', 'total_bedrooms', 'population', 'households', 'median_income', 'ocean_proximity']
[[-122.23  37.88  41.0 ... 126.0  8.3252 'NEAR BAY']
 [-122.22  37.86  21.0 ... 1138.0  8.3014 'NEAR BAY']
 [-122.24  37.85  52.0 ... 177.0  7.2574 'NEAR BAY']
 ...
 [-121.22  39.43  17.0 ... 433.0  1.7 'INLAND']
 [-121.32  39.43  18.0 ... 349.0  1.8672 'INLAND']
 [-121.24  39.37  16.0 ... 530.0  2.3886 'INLAND']]
(19648, 9)
```

```
In [13]: 1 # Extracting the target features from the data
2 y = housing["median_house_value"].copy()
3 y = y.values
4
5 print(y) # Or display(y_train)
6 print(y.shape)

[452600. 358500. 352100. ... 92300. 84700. 89400.]
(19648,)
```

Data Cleaning

At this stage in our workflow we may want to deal with duplicated/missing values and (optionally) fix/remove outliers. Let's start with the former.

We want to remove duplicates as they may bias our fitted model. In other words, we may potentially *overfit* to this subset of points. However, care should usually be taken to check they are not *real* data with identical values.

There a number of ways we could identify duplicates, the simplest (and the approach we'll focus on) is just to find observations with the same feature values. Of course this will not identify things such as spelling errors, missing values, address changes, use of aliases, etc. For those such things, more complicated methods along with manual assessment is needed.

► Exercise 8 (CORE)

Are there any duplicated values in the data? If so how many?

Hint

```
In [14]: 1 housing.duplicated().sum()
```

```
Out[14]: 0
```

⚡ Now, is a good point to switch driver and navigator

Outlier Detection

An **Outlier** is a data point that lies abnormally far from other observations and may distort the model fit and results. In general, it is a good idea to examine if any outliers are present during preprocessing. However, it may not always be appropriate to drop these observations without explicit knowledge and testing if they are really outliers or not.

We will use basic statistics in order to try to identify outliers. A simple method of detecting outliers is to use the **inter-quartile range (IQR) proximity rule** (Tukey fences) which states that a value is an outlier if it falls outside these boundaries:

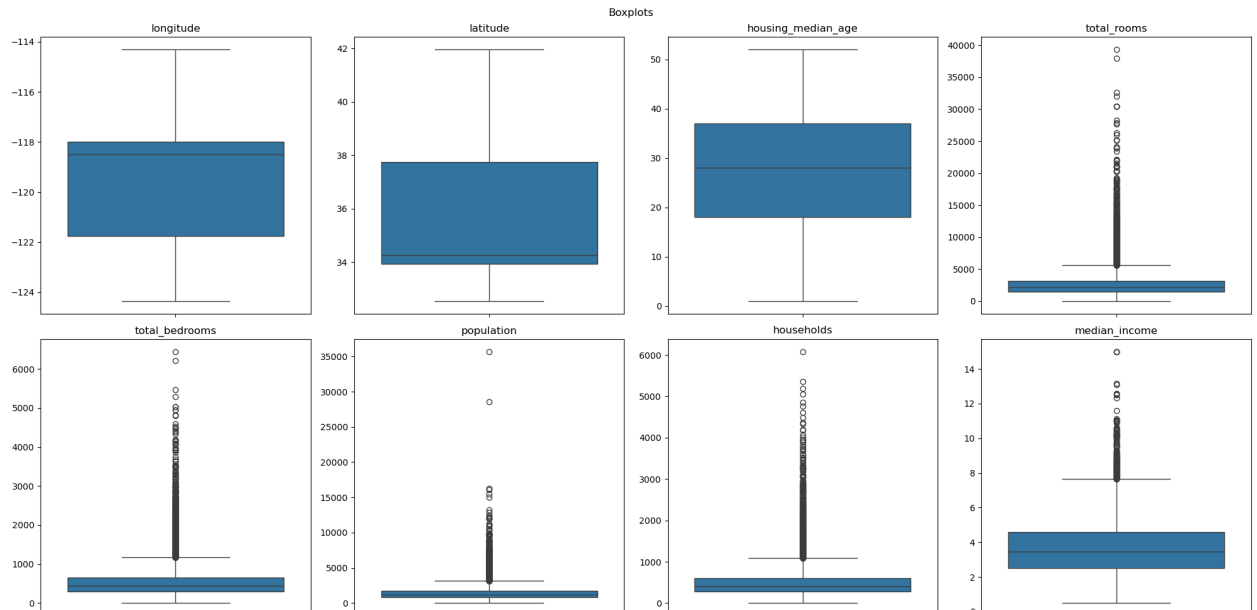
- Upper boundary = 75th quantile + (IQR * k)
- Lower boundary = 25th quantile - (IQR * k)

where IQR = 75th quantile - 25th quantile (the length of the box in the box plot). Here, k is a nonnegative constant which is typically set to 1.5. However, it is also common practice to find extreme values by setting k to 3.

► Exercise 9 (EXTRA)

Can you identify any potential outliers using the generated boxplots below ?

```
In [15]: 1 fig, axes = plt.subplots(figsize = (20,10), ncols = X.shape[-1]//2, nrows = 2, sharex = True)
2 axes = axes.flatten()
3
4 for i, ax in enumerate(axes):
5     sns.boxplot(y = X[:,i], ax = ax)
6     ax.set_title(features[i])
7     ax.set_ylabel("")
8
9 plt.suptitle("Boxplots")
10 plt.tight_layout()
11 plt.show()
```



The boxplots in the provided image show several outliers in the housing dataset across multiple features:

Longitude: There are no apparent outliers in longitude values as the data appears to be tightly clustered within the box plot range.

Latitude: The latitude values show no outliers, as the data distribution is contained within the box plot range.

Housing Median Age: No significant outliers are visible in the housing median age variable. The data is well-distributed within the interquartile range.

Total Rooms: This feature shows a significant number of outliers, with values extending well beyond the upper whisker. These values may indicate extremely large housing units or highly variable room counts that differ significantly from the bulk of the data.

Total Bedrooms: Outliers are observed at the higher end of the range, with values exceeding 4,000 bedrooms. These could be indicative of exceptionally large homes or unusual data entries.

Population: The population variable displays several extreme outliers, especially values close to 35,000, suggesting some locations with disproportionately large populations compared to the general distribution.

Households: A significant number of outliers are found at the higher end, with values approaching or exceeding 6,000 households. These may represent large urban areas or regions with high-density housing.

Median Income: There are also a considerable number of outliers in median income, with values well above 10. These outliers likely reflect areas with very high median incomes compared to the general population.

These outliers might represent unusual areas or data entry errors, and further investigation could be necessary to determine their cause and handle them appropriately for analysis.

- After detecting outliers we may want to set them to NA values so we can use subsequent methods, e.g to remove them (NA Removal).
- In general, it is a good idea to make our methods into functions that are compatible with scikit-learn, so that we can use them in a Pipeline later on.
- The removal of outliers is NOT an easy task, should not be applied all the time, especially if you are interested in rare events. In general, without any specific knowledge for the problem at hand or without statistical examination of outliers via suitable testing tools, it is not recommended to remove these observations by looking only at the above boxplots.
- Another issue might be related to the sample size of your data set. Even though the `housing` data is large, in some cases, when you drop those observations you can discard important information unwittingly!

For the sake of simplicity, we are keeping all observations at that moment for this workshop.

Missing Data

Most ML models cannot handle missing values, and as we saw earlier, there are some present in `total_bedrooms`.

There are a number of ways we can deal with missing values but to start with, let's just **remove NA values**. We can do this in two ways by either:

1. Getting rid of the corresponding observations (deleting the corresponding rows).
2. Getting rid of the whole attribute (deleting the corresponding columns).

To do this outside of a pipeline is relatively straight forward, we just run `housing.dropna()` with either the `axis` set to 0 or 1 (depending if we want to remove rows or columns) before splitting our data into features (`X`) and outputs (`y`).

However if we want to do this in a pipeline we can use the *transformer* `DropMissingData` (https://feature-engine.trainindata.com/en/1.8.x/api_doc/imputation/DropMissingData.html) from `feature_engine.imputation`. Instead of simply dropping missing data, we may instead want to use other imputation methods are available in `sklearn.impute` (<https://scikit-learn.org/1.5/api/sklearn.impute.html>).

Transformers

If we want to **alter the features** of our data using a function that is compatible with a `Pipeline`, we need to a *transformer*.

- Transformers are classes that follow the scikit-learn API in Scikit-Learn [clean](https://scikit-learn.org/stable/modules/preprocessing.html#preprocessing) (<https://scikit-learn.org/stable/modules/preprocessing.html#preprocessing>), [impute](https://scikit-learn.org/1.5/modules/impute.html) (<https://scikit-learn.org/1.5/modules/impute.html>), [reduce](https://scikit-learn.org/stable/modules/unsupervised_reduction.html#data-reduction) (https://scikit-learn.org/stable/modules/unsupervised_reduction.html#data-reduction), [expand](https://scikit-learn.org/stable/modules/kernel_approximation.html#kernel-approximation) (https://scikit-learn.org/stable/modules/kernel_approximation.html#kernel-approximation), or [generate](https://scikit-learn.org/stable/modules/feature_extraction.html#feature-extraction) (https://scikit-learn.org/stable/modules/feature_extraction.html#feature-extraction) feature representations.
- Transformers are classes with a `.fit()` method, which learn model parameters (e.g. mean for mean imputation) from a training set, and a `.transform()` method which applies this transformation model to data. To create a custom transformer, all you need is to create a class that implements three methods: `fit()`, `transform()`, and `fit_transform()`.

Therefore to transform a dataset, each sampler implements:

```
obj.fit(data)
data_transformed = obj.transform(data)
```

or simply...

```
data_transformed = obj.fit_transform(data)`
```

See more details: https://scikit-learn.org/stable/data_transforms.html (https://scikit-learn.org/stable/data_transforms.html). In the following subsections, we will see examples of *transformers* for categorical and numerical variables.

► Exercise 10 (CORE)

Use `dropMissingData` to remove the missing observations from `X`. What is the shape after dropping the missing observations?

Notes

- It may be tempting to overwrite `X` while working on our pre-processing steps. **Don't do this!** We will run these objects through these steps inside our pipeline later so if you want to test your function make sure to assign the output to temporary objects (e.g. `X_`).

```
In [16]: 1 from feature_engine.imputation import DropMissingData
2 imputer = DropMissingData()
3 housing_cleaned = imputer.fit_transform(housing)
4
```

```
-----
AttributeError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_11628\3821007164.py in <module>
      1 from feature_engine.imputation import DropMissingData
      2 imputer = DropMissingData()
----> 3 housing_cleaned = imputer.fit_transform(housing)

~\anaconda3\lib\site-packages\sklearn\utils\_set_output.py in wrapped(self, X, *args, **kwargs)
    317 @wraps(f)
    318 def wrapped(self, X, *args, **kwargs):
--> 319     data_to_wrap = f(self, X, *args, **kwargs)
    320     if isinstance(data_to_wrap, tuple):
    321         # only wrap the first output for cross decomposition

~\anaconda3\lib\site-packages\sklearn\base.py in fit_transform(self, X, y, **fit_params)
    916 if y is None:
    917     # fit method of arity 1 (unsupervised transformation)
--> 918     return self.fit(X, **fit_params).transform(X)
    919 else:
    920     # fit method of arity 2 (supervised transformation)

~\anaconda3\lib\site-packages\sklearn\utils\_set_output.py in wrapped(self, X, *args, **kwargs)
    317 @wraps(f)
    318 def wrapped(self, X, *args, **kwargs):
--> 319     data_to_wrap = f(self, X, *args, **kwargs)
    320     if isinstance(data_to_wrap, tuple):
    321         # only wrap the first output for cross decomposition

~\anaconda3\lib\site-packages\feature_engine\imputation\drop_missing_data.py in transform(self, X)
    181 """
    182
--> 183     X = self._transform(X)
    184
    185     if self.threshold:

~\anaconda3\lib\site-packages\feature_engine\imputation\base_imputer.py in _transform(self, X)
     30 """
     31     # Check method fit has been called
----> 32     check_is_fitted(self)
     33
     34     # check that input is a dataframe

~\anaconda3\lib\site-packages\sklearn\utils\validation.py in check_is_fitted(estimator, attributes, msg, all_or_any)
    1749     raise TypeError("%s is not an estimator instance." % (estimator))
    1750
-> 1751     tags = get_tags(estimator)
    1752
    1753     if not tags.requires_fit and attributes is None:

~\anaconda3\lib\site-packages\sklearn\utils\_tags.py in get_tags(estimator)
    428     for klass in reversed(type(estimator).mro()):
    429         if "__sklearn_tags__" in vars(klass):
--> 430             sklearn_tags_provider[klass] = klass.__sklearn_tags__(estimator) # type: ignore[attr-de
fined]
    431             class_order.append(klass)
    432             elif "__more_tags__" in vars(klass):

~\anaconda3\lib\site-packages\sklearn\base.py in __sklearn_tags__(self)
    857
    858     def __sklearn_tags__(self):
--> 859     tags = super().__sklearn_tags__()
    860     tags.transformer_tags = TransformerTags()
    861     return tags

AttributeError: 'super' object has no attribute '__sklearn_tags__'
```

Note that later in the course, we will also learn about **Resamplers** which allow us to *alter the number observations* in our data using a function that is compatible with a Pipeline . For more details: <https://imbalanced-learn.org/stable/introduction.html> (<https://imbalanced-learn.org/stable/introduction.html>)

Data Imputation

Instead of removing the missing data we can set it to some value. To do this, Scikit-Learn gives us a handy `SimpleImputer` (<https://scikit-learn.org/1.5/modules/generated/sklearn.impute.SimpleImputer.html>) method which provides simple strategies (e.g. "mean", "median" for numerical features and "most_frequent" for categorical features). You can also add a missing indicator, which may be useful in the case when missing features may provide information for predicting the target (e.g. obese patients may prefer not to report bmi, thus, this missingness could be useful for estimating the risk of health conditions or diseases). Beyond simple imputation strategies, sklearn also provides more advanced imputation strategies in `IterativeImputer` (<https://scikit-learn.org/1.5/modules/generated/sklearn.impute.IterativeImputer.html#sklearn.impute.IterativeImputer>) and `KNNImputer` (<https://scikit-learn.org/1.5/modules/generated/sklearn.impute.KNNImputer.html>). Other strategies are also available in `SimpleImputer`.

```
In [19]: 1 from sklearn.impute import SimpleImputer
2 num_imputer = SimpleImputer(strategy="median")
3 X = housing.copy()
4 X_imp = num_imputer.fit_transform(X)

-----
ValueError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_11628\1860213672.py in <module>
      2 num_imputer = SimpleImputer(strategy="median")
      3 X = housing.copy()
----> 4 X_imp = num_imputer.fit_transform(X)

~\anaconda3\lib\site-packages\sklearn\utils\_set_output.py in wrapped(self, X, *args, **kwargs)
    317 @wraps(f)
    318 def wrapped(self, X, *args, **kwargs):
--> 319     data_to_wrap = f(self, X, *args, **kwargs)
    320     if isinstance(data_to_wrap, tuple):
    321         # only wrap the first output for cross decomposition

~\anaconda3\lib\site-packages\sklearn\base.py in fit_transform(self, X, y, **fit_params)
    916 if y is None:
    917     # fit method of arity 1 (unsupervised transformation)
--> 918     return self.fit(X, **fit_params).transform(X)
    919 else:
    920     # fit method of arity 2 (supervised transformation)

~\anaconda3\lib\site-packages\sklearn\base.py in wrapper(estimator, *args, **kwargs)
    1387 )
    1388 ):
--> 1389     return fit_method(estimator, *args, **kwargs)
    1390
    1391 return wrapper

~\anaconda3\lib\site-packages\sklearn\impute\_base.py in fit(self, X, y)
    432     Fitted estimator.
    433     """
--> 434     X = self._validate_input(X, in_fit=True)
    435
    436     # default fill_value is 0 for numerical input and "missing_value"

~\anaconda3\lib\site-packages\sklearn\impute\_base.py in _validate_input(self, X, in_fit)
    359 )
    360 )
--> 361     raise new_ve from None
    362 else:
    363     raise ve

ValueError: Cannot use median strategy with non-numeric data:
could not convert string to float: 'NEAR BAY'
```

When we applied this to our data, we get the following error:

```
ValueError: Cannot use median strategy with non-numeric data:
could not convert string to float:
```

This is because the "median" strategy can only be used with numerical attributes so we need a way of only applying imputation to certain attributes. As with most things in Python, there are a number of options available to us (you may have already used one of them in answers to previous questions):

- We could temporarily remove the categorical feature from our data to apply our function

```
X_num = X[:, :-1]
X_num = num_imputer.fit_transform(X_num)
```

- Apply the function to a subset of the data and assign the output to the same subset

```
X[:, :-1] = num_imputer.fit_transform(X[:, :-1])
```

However scikit-learn has a handy function to specify what column we want to apply a function to (its also compatible with Pipeline s as well).

```
In [20]: 1 from sklearn.compose import ColumnTransformer
2 features = X.columns
3 num_cols_imputer = ColumnTransformer(
4     # apply the `num_imputer` to all columns apart from the last
5     [ ("num", num_imputer, list(range(X.shape[1] - 1))) ],
6     # don't touch all other columns, instead concatenate it on the end of the
7     # changed data.
8     remainder = "passthrough"
9 )
10
11 X_ = num_cols_imputer.fit_transform(X)
12 print("Number of Missing Values")
13 pd.DataFrame(X_, columns = features).isna().sum()
```

Number of Missing Values

```
Out[20]: longitude      0
latitude      0
housing_median_age    0
total_rooms         0
total_bedrooms       0
population          0
households          0
median_income        0
median_house_value   0
ocean_proximity     0
dtype: int64
```

```
In [21]: 1 pd.DataFrame(X_, columns = features)
```

```
Out[21]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	452600.0
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	358500.0
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	352100.0
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	341300.0
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	342200.0
...
19643	-121.09	39.48	25.0	1665.0	374.0	845.0	330.0	1.5603	78100.0
19644	-121.21	39.49	18.0	697.0	150.0	356.0	114.0	2.5568	77100.0
19645	-121.22	39.43	17.0	2254.0	485.0	1007.0	433.0	1.7	92300.0
19646	-121.32	39.43	18.0	1860.0	409.0	741.0	349.0	1.8672	84700.0
19647	-121.24	39.37	16.0	2785.0	616.0	1387.0	530.0	2.3886	89400.0

19648 rows × 10 columns

After using `.fit` we can see that this method stores the median value for each attribute on the training set. This value can be used when validating and testing the model as it will be used if there is missing data in the new data.

Note

- To access the fitted `num_imputer` in this case, you first needed to subset the `transformers_` in `ColumnTransformer` to only the first tuple (`[0]`), and then get the second object (`[1]`) in the list (the `num_imputer`).

```
In [22]: 1 display(num_cols_imputer.transformers_[0][1].statistics_)

array([-1.18500e+02,  3.42700e+01,  2.80000e+01,  2.11100e+03,
        4.36000e+02,  1.17900e+03,  4.11000e+02,  3.44905e+00,
        1.73600e+05])
```

⚠ Now, is a good point to switch driver and navigator

Feature Engineering

As discussed in the lectures, feature engineering is where we extract features from data and transform them into formats that are suitable for machine learning models. Today, we will have a look at two main cases that are present in our data: **categorical** and **numerical** values.

Feature engineering also requires a *transformer* class to **alter the features**.

Categorical Variables

- In the dataset, we have a text attribute (`ocean_proximity`) that we already had to treat differently when cleaning the data. This extends to feature engineering as well, where we need to use separate methods than those used with numerical variables.
- If we look at the unique values of this attribute, we will see that there are a limited number of possible values which represent a category. We need a way of encoding this information into our modeling framework **by converting our string/categorical variable into a numeric representation** that can be included in our models.

If we have a binary categorical variable (two levels) we could do this by picking one of the categorical levels and encode it as 1 and the other level as 0.

However, in this case as we have multiple categories, we would probably want to use another encoding method. To illustrate, we can try encoding the categorical feature `ocean_proximity` using both the `OrdinalEncoder` and `OneHotEncoder` available in `sklearn.preprocessing`.

Side Notes

- The output of the `OneHotEncoder` provided in Scikit-Learn is a SciPy *sparse matrix*, instead of a NumPy array. These are useful when you have lots of categories as your matrix becomes mostly full of 0's. To store all these 0's takes up unnecessary memory, so instead a sparse matrix just stores the location of nonzero elements. The good news is that you can use a sparse matrix similar to a numpy matrix, but if you wanted to, you can convert it to a dense numpy matrix using `.toarray()`.
- The above does not seem to be the case if passed through a `ColumnTransformer`.

```
In [23]: 1 from sklearn.preprocessing import OrdinalEncoder
2 # Defining the OrdinalEncoder
3 ordinal_encoder = OrdinalEncoder()
4
5 encoder = ColumnTransformer([
6     # pass through all columns with numerical values
7     ("num", "passthrough", list(range(X.shape[1] - 1))),
8     # apply the ordinal_encoder to the last column
9     ("cat", ordinal_encoder, [X.shape[1] - 1]),
10 ])
11
12 # fitting the encoder defined above
13 housing_cat_encoded = encoder.fit_transform(X)
14
15 display(dict(zip(list(encoder.transformers_[1][1].categories_[0]), range(5))))
```

```
{'<1H OCEAN': 0, 'INLAND': 1, 'ISLAND': 2, 'NEAR BAY': 3, 'NEAR OCEAN': 4}
```

```
In [24]: 1 from sklearn.preprocessing import OneHotEncoder
2 # Defining the OneHotEncoder
3 onehot_encoder = OneHotEncoder()
4
5 encoder = ColumnTransformer([
6     # pass through all columns with numerical values
7     ("num", "passthrough", list(range(X.shape[1] - 1))),
8     # apply the onehot_encoder to the last column
9     ("cat", onehot_encoder, [X.shape[1] - 1]),
10 ])
11
12 housing_cat_1hot = encoder.fit_transform(X)
13 cats = pd.Series(X.iloc[:, -1]).unique()
14 pd.DataFrame(encoder.transformers_[1][1].transform(cats.reshape(-1, 1)).toarray(), columns = cats).astype(int)
```

```
C:\Users\TRISHA\anaconda3\lib\site-packages\sklearn\utils\validation.py:2739: UserWarning: X does not have valid feature names, but OneHotEncoder was fitted with feature names
warnings.warn(
```

Out[24]:

	NEAR BAY	<1H OCEAN	INLAND	NEAR OCEAN	ISLAND
0	0	0	0	1	0
1	1	0	0	0	0
2	0	1	0	0	0
3	0	0	0	0	1
4	0	0	1	0	0

► Exercise 11 (CORE)

- What is the main difference between two methods regarding the obtained features ?
- Which encoding method do you think is most appropriate for this variable and why?

1

Another handy feature of `OneHotEncoder` and `OrdinalEncoder` is that infrequent categories can be aggregated into a single feature/value. The parameters to enable the gathering of infrequent categories are `min_frequency` and `max_categories`.

When there are many unordered categories, another useful encoding scheme is `TargetEncoder` (<https://scikit-learn.org/1.5/modules/generated/sklearn.preprocessing.TargetEncoder.html#sklearn.preprocessing.TargetEncoder>) which uses the target mean conditioned on the categorical feature for encoding unordered categories. Whereas one-hot encoding would greatly inflate the feature space if there are a very large number of categories (e.g. zip code or region), `TargetEncoder` is more parsimonious.

► Exercise 12 (EXTRA)

Use the `max_categories` attribute to set the maximum number of categories to 4. Use the `get_feature_names_out()` method of `OneHotEncoder` to print the new category names. Which two features have been combined?

In []:

1

Type *Markdown* and LaTeX: α^2

Numerical Variables

Box-Cox Transformation

In some cases, we may wish to apply transformations to our data, so that they have a more Gaussian distribution. For example, log transformations are useful for altering count data to have a more normal distribution as they pull in the more extreme high values relative to the median, while stretching back extreme low values away from the median. You can use a log transformation with either the pre-made `LogTransformer()` from `feature_engine.transformation`, or a custom function and `sklearn.preprocessing.FunctionTransformer`.

More generally, the natural logarithm, square root, and inverse transformations are special cases of the **Box-Cox** family of transformations (Box and Cox 1964). The question is why do we need such a transformation and when?

- Note that, the method is typically used to transform the outcome, but can also be used to transform predictors.
- The method assumes that the variable takes only positive values. If there are any zero or negative values, a modification of the Box-Cox that allows non-positive values (Box-Cox with negatives, BCN) can be used (Hawkins and Weisberg 2017).
- If there are zero values, but no negative values, an option is to add a small number (e.g. 1) before carrying out the transformation, although if there are a large number of zero values then it is possible no Box-Cox transformation will sufficiently normalize the variable.
- In general, transformations can make interpretations more difficult, thus you should think carefully if they are needed, particularly if they only result in modest improvements in model performance. Moreover, finding a suitable transformation is typically a trial-and-error process.

Originally, the Box-Cox procedure was developed to transform the outcome and uses maximum likelihood estimation to estimate a transformation parameter λ .

$$\tilde{y} = \begin{cases} \frac{y^\lambda - 1}{\lambda}, & \lambda \neq 0 \\ \log y, & \lambda = 0 \end{cases}$$

- λ is estimated from the data.
- Because the parameter of interest is in the exponent, this type of transformation is called a power transformation
- Some values of λ map to common transformations, such as;

- $\lambda = 1$ (no transformation)
- $\lambda = 0$ (log)
- $\lambda = 0.5$ (square root)
- $\lambda = -1$ (inverse)

- Using the code below, if `lmbda=None` then the function will "find the lambda that maximizes the log-likelihood function and return it as the second output argument"
- Notice that we can not use `lambda` directly since it conflicts with the available object called `lambda`, this is the reason we preferred the indicator name as `lmbda`


```
In [ ]: 1 fig, axes = plt.subplots(figsize = (15,5), ncols = 4, nrows=2, sharey = True)
2 axes = axes.flatten()
3 sns.histplot(data = X[:,6], ax = axes[0])
4 axes[0].set_title("Raw Counts")
5
6 for i, lmbda in enumerate([0, 0.25, 0.5, 0.75, 1., 1.25, 1.5]):
7
8     house_box_ = stats.boxcox(X[:, 6].astype(float), lmbda = lmbda)
9     sns.histplot(data = house_box_, ax = axes[i + 1])
10    axes[i + 1].set_title("$\lambda$ = {}".format(lmbda))
11
12 plt.tight_layout()
13 plt.show()
```

► Exercise 13 (EXTRA)

- For the variable `households`, based on the `boxcox` transform shown above, which values of λ may be useful?
- Apply a similar code snippet to `median_house_value`. Which values of λ may be useful?

Type *Markdown* and LaTeX: α^2

In []:

1

Type *Markdown* and LaTeX: α^2

Other feature types

Feature engineering for other feature types beyond numerical categorical are also available in sklearn (e.g. for [text and images](https://scikit-learn.org/1.5/api/sklearn.feature_extraction.html) (https://scikit-learn.org/1.5/api/sklearn.feature_extraction.html)) and feature engine (e.g. for [Datetime](https://feature-engine.trainindata.com/en/1.8.x/user_guide/datetime/index.html) (https://feature-engine.trainindata.com/en/1.8.x/user_guide/datetime/index.html) and for [time series](https://feature-engine.trainindata.com/en/1.8.x/user_guide/timeseries/index.html) (https://feature-engine.trainindata.com/en/1.8.x/user_guide/timeseries/index.html)).

Feature Combinations

- Looking at the data's attributes we may also want to manually combine them into features that are either a more meaningful representation of the data or have better properties.
- For example, we know **the number of rooms** in a district, but this may be more useful to combine with the **number of households** so that we have **a measure of rooms per household**.

```
In [ ]: 1 col_names = "total_rooms", "households"
2 rooms_ix, households_ix = [housing.columns.get_loc(c) for c in col_names]
3
4 rooms_per_household = X[:, rooms_ix] / X[:, households_ix]
5 rooms_per_household
```

► Exercise 14 (EXTRA)

- Can you think of other combinations that may be useful?
- Create a custom transformer that creates these new combinations of features using the `FunctionTransformer`.

Hint

Type *Markdown* and LaTeX: α^2

In []:

1

⌘ Now, is a good point to switch driver and navigator

Feature Scaling

As we will discuss in later weeks, many machine learning algorithms are sensitive to the scale and magnitude of the features, and especially differences in scales across features. For these algorithms, feature scaling will improve performance.

For example, let's investigate the range of the features in our dataset:

```
In [ ]: 1 fig, ax = plt.subplots(figsize=(15,5))
2
3 plt.boxplot(X[:, :-1], vert = False)
4 plt.xscale("symlog")
5 plt.ylabel("Feature")
6 plt.xlabel("Feature Range")
7
8 ax.set_yticklabels(features[:, :-1])
9
10 plt.suptitle("Feature Range for the Training Set")
11 plt.tight_layout()
12 plt.show()
```

There are various options in scikit learn for feature scaling:

- Standardization (`preprocessing.StandardScaler`)
- Min-Max Scaling (`preprocessing.MinMaxScaler`)
- l2 Normalization (`preprocessing.normalize`)
- RobustScaler(`preprocessing.RobustScaler`)
- Scale with maximum absolute value (`preprocessing.MaxAbsScaler`)

- As scaling generally improves the performance of most models when features cover a range of scales, it is probably a good idea to apply some sort of scaling to our data before fitting a model.
- *Standardization* (or *variance scaling*), is the most common, but there are a number of other types, as listed above.

► Exercise 15 (EXTRA)

Try implementing different scalers for the `total_rooms` variable to see the main differences.

```
In [ ]:
```

```
1
```

Combining into a Pipeline

Due to the focus of this week, let's look at combining different feature engineering steps together to make different model pipelines.

- We want to create a basic model pipeline that treats the **numerical** and **categorical** attributes differently.
- To achieve this we are going to use a combination of `Pipeline` and `ColumnTransformer` objects.
- We also need to supply the pipeline with a model. For now let's use a linear regression model, which we will learn in more details in week 4.

We can of course do this outside of a pipeline, but remember to be careful about data leakage.

```

In [ ]: 1 from sklearn.preprocessing import StandardScaler
2 from sklearn.linear_model import LinearRegression
3 from sklearn.pipeline import Pipeline
4 from sklearn.preprocessing import OneHotEncoder
5 from sklearn import set_config
6 set_config(display='diagram')
7
8 numcols = list(range(X.shape[1]-1))
9 catcols = [X.shape[1]-1]
10
11 num_pre = Pipeline([
12     ("num_impute", SimpleImputer(strategy="median")),
13     ("num_scale", StandardScaler())])
14
15 cat_pre = Pipeline([
16     ("cat_impute", SimpleImputer(strategy="constant")),
17     ("cat_encode", OneHotEncoder(drop='first'))])
18
19 reg_pipe_1 = Pipeline([
20     ("pre_processing", ColumnTransformer([("num_pre", num_pre, numcols),
21                                         ("cat_pre", cat_pre, catcols)])),
22     ("model", LinearRegression())
23 ])
24
25 # Alternative and equivalent model avoiding nested pipelines
26 # reg_pipe_1 = Pipeline([
27 #     ("impute", ColumnTransformer([("num_imp", SimpleImputer(strategy="median"), numcols),
28 #                                   ("cat_imp", SimpleImputer(strategy="constant"), catcols)])),
29 #     ("transform", ColumnTransformer([("num_trns", StandardScaler(), numcols),
30 #                                       ("cat_trns", OneHotEncoder(drop='first'), catcols)])),
31 #     ("model", LinearRegression())
32 # ])
33
34 display(reg_pipe_1)

```

```

In [ ]: 1 reg_pipe_1.fit(X,y)
2 # Print the R squared (ranges 0 to 1, with higher values better)
3 print(round(reg_pipe_1.score(X, y), 3))

```

```

In [ ]: 1 # Print the coefficients
2 reg_pipe_1[1].coef_
3 # reg_pipe_1['model'].coef_

```

Let's try some other combinations of the pre-processing and feature engineering steps that we have learned about this week.

```

In [ ]: 1 # NOTE: ADDITIONAL transformer for the pipeline
2 # Creating the Transformer for box-cox implementation
3 from sklearn.preprocessing import FunctionTransformer
4
5 def boxcoxarray(X, y=None, lmbda=None):
6     X_ = X.copy()
7     for i in range(X.shape[1]):
8         if lmbda == None:
9             X[:,i], bc_params_ = stats.boxcox(X[:,i], lmbda = lmbda)
10        else:
11            X[:,i] = stats.boxcox(X[:,i], lmbda = lmbda)
12    return X_
13
14 # Creating our transformer from the pre-defined boxcoxarray callable function:
15 boxtransformer = FunctionTransformer(boxcoxarray, validate = True)

```

```
In [ ]: 1 # Reg Pipe 2
2
3 # Define column indices
4 numcols = [0, 1, 2, 7]
5 countcols = [3, 4, 5, 6]
6 catcols = [8]
7
8 # Reg Pipe 2
9 num_pre = Pipeline([
10     ("num_impute", SimpleImputer(strategy="median")),
11     ("num_scale", StandardScaler())])
12
13 count_pre = Pipeline([
14     ("num_impute", SimpleImputer(strategy="median")),
15     ("num_transform", boxtransformer),
16     ("num_scale", StandardScaler())])
17
18 cat_pre = Pipeline([
19     ("cat_impute", SimpleImputer(strategy="constant")),
20     ("cat_encode", OneHotEncoder(drop='first'))])
21
22 # Overall ML pipeline including all
23 reg_pipe_2 = Pipeline([
24     ("pre_processing", ColumnTransformer([
25         ("num_pre", num_pre, numcols), # Applied to the columns/variables in 0, 1, 2, 7 index
26         ("count_pre", count_pre, countcols), # Applied to the columns/variables in 3, 4, 5, 6 index
27         ("cat_pre", cat_pre, catcols)])), # Applied to the columns/variables in 8 index only (ocean_proximity
28     ("model", LinearRegression())
29 ])
30
31
32 display(reg_pipe_2)
```

```
In [ ]: 1 reg_pipe_2.fit(X,y)
2 # Print the R squared (ranges 0 to 1, with higher values better)
3 print(round(reg_pipe_2.score(X, y), 3))
```

```
In [ ]: 1 # Print the coefficients
2 reg_pipe_2[1].coef_
```

```
In [ ]: 1 # Reg Pipe 3
2 from feature_engine.transformation import LogTransformer
3
4 numcols = [0, 1]
5 countcols = [2, 3, 4, 5, 6, 7]
6 catcols = [8]
7
8 num_pre = Pipeline([
9     ("num_impute", SimpleImputer(strategy="median")),
10    ("num_scale", StandardScaler())])
11
12 num2_pre = Pipeline([
13    ("num_impute", SimpleImputer(strategy="median")),
14    ("num_transform", LogTransformer()),
15    ("num_scale", StandardScaler())])
16
17 cat_pre = Pipeline([
18    ("cat_impute", SimpleImputer(strategy="constant")),
19    ("cat_encode", OneHotEncoder(drop='first'))])
20
21 # Overall ML pipeline including all
22 reg_pipe_3 = Pipeline([
23    ("pre_processing", ColumnTransformer([
24        ("num_pre", num_pre, numcols), # Applied to the columns/variables in 0, 1 index
25        ("num2_pre", count_pre, countcols), # Applied to the columns/variables in 2, 3, 4, 5, 6, 7 index
26        ("cat_pre", cat_pre, catcols)])), # Applied to the columns/variables in 8 index only (ocean_proximity
27    ("model", LinearRegression())
28 ])
29
30 display(reg_pipe_3)
```

```
In [ ]: 1 reg_pipe_3.fit(X,y)
2 # Print the R squared (ranges 0 to 1, with higher values better)
3 print(round(reg_pipe_3.score(X, y), 3))
```

```
In [ ]: 1 # Print the coefficients
2 reg_pipe_3[1].coef_c
```

► Exercise 16 (CORE)

Explain in words what are the differences in pre-processing and feature engineering steps in the three models above.

Type *Markdown* and LaTeX: α^2

► Exercise 17 (EXTRA)

Try to create your own pipeline by modifying at least one of the pre-processing and feature engineering steps above. What have you decided to change and why?

Summary

This week we covered a lot of ground!

- We've looked at some methods for pre-processing our data, cleaning and preparing it, as well as how to engineer some features.
- This is by **no means a complete collection of all the methods available** as covering more would go beyond the scope of this course (for those interested in learning more, have a look though the given companion readings).
- For example, we did not touch on handling text and dates/time much. These topics are quite complex and have enough materials to cover their own courses.

Competing the Worksheet

At this point you have hopefully been able to complete all the CORE exercises and attempted the EXTRA ones. Now is a good time to check the reproducibility of this document by restarting the notebook's kernel and rerunning all cells in order.

Before generating the PDF, please go to Edit -> Edit Notebook Metadata and change 'Student 1' and 'Student 2' in the **name** attribute to include your name. If you are not able to edit the Metadata, then please add a markdown cell at the beginning with the names of all group members.

Once that is done and you are happy with everything, you can then run the following cell to generate your PDF.

```
In [ ]: 1 !jupyter nbconvert --to pdf mlp_week01.ipynb
```

Once generated, please submit this PDF on Learn page by 16:00 PM on the Friday of the week the workshop was given. Note that:

- You don't need to finish everything, but you should have had a substantial attempt at the bulk of the material, particularly the CORE tasks.
- If you are having trouble generating the pdf, please ask a tutor or post on piazza.
- As a back option, you can simply select the print option in your browser which should generate a high-quality pdf of your complete notebook. Basically, right click → Click Print → Choose 'Save as PDF' option instead of default printing destination (Check that your all materials are visible in the print preview) → Click 'Save' to get your pdf file.